

Advanced Java programming (Java EE)

Java Persistence API

Based on the JPA presentation from javabeat.net

Topics

- **Introduction to java persistence**
 - The Java Persistence API
 - Entities
 - EntityManager & the Persistent Context
 - Persistence Units
 - Exceptions
 - JPA Query Language
-

Introduction

- Previously we learnt about
 - JDBC
 - Data Access Objects (DAO) and Data Transfer Objects (DTO)
 - 1. In JDBC, we "hard coded" SQL into our application
 - 2. Then used Data Source/Connection Pooling
 - 3. Then used DAO/DTO
 - 4. But this just "hides" implementation from our business logic, you still implement DAO with JDBC
-

Issues not solved

- However,
 - We still have to understand a lot of implementation details (eg: connections, statements, resultsets etc)
 - What about relationships? Joins? Inheritance?
 - Object \leftrightarrow database impedance mismatch
 - J2EE tried to solve this with
"Entity Enterprise JavaBeans (EJB)"
 - Simpler alternatives included
Object Relational Mapping (ORM) tools:
 - e.g. Java Data Objects (JDO), Hibernate, iBatis, TopLink
-

Java EE to the rescue

- Java SE 5 added new constructs to Java language
 - Generics
 - Annotations
 - Enumerations
 - Java EE 5 used these features to provide
 - Ease of development
 - "Dependency injection"
 - Meaningful defaults, "code by exception"
 - Simplified EJB
 - New Java Persistence API (JPA) replaced Entity EJB
-

Java EE 5 persistence

- Java EE 5 still keeps JDBC
 - EJB 3 spec (JSR 220) split into 2:
 1. Session Beans, Message Beans
 2. Java Persistence API (JPA)
 - JPA jointly developed by TopLink, Hibernate, JDO, EJB vendors and individuals
 - JPA can also be used in Java SE 5 without a container!!!!
-

- Java Persistence consists of three areas:
 - The Java Persistence API
 - The query language
 - Object/relational mapping metadata
 - JPA implementation
 - Reference implementation: TopLink (GlassFish project)
 - Most ORM vendors now have JPA interface
 - eg: Hibernate-JPA, EclipseLink (based on TopLink), OpenJPA (based on BEA Kodo)
-

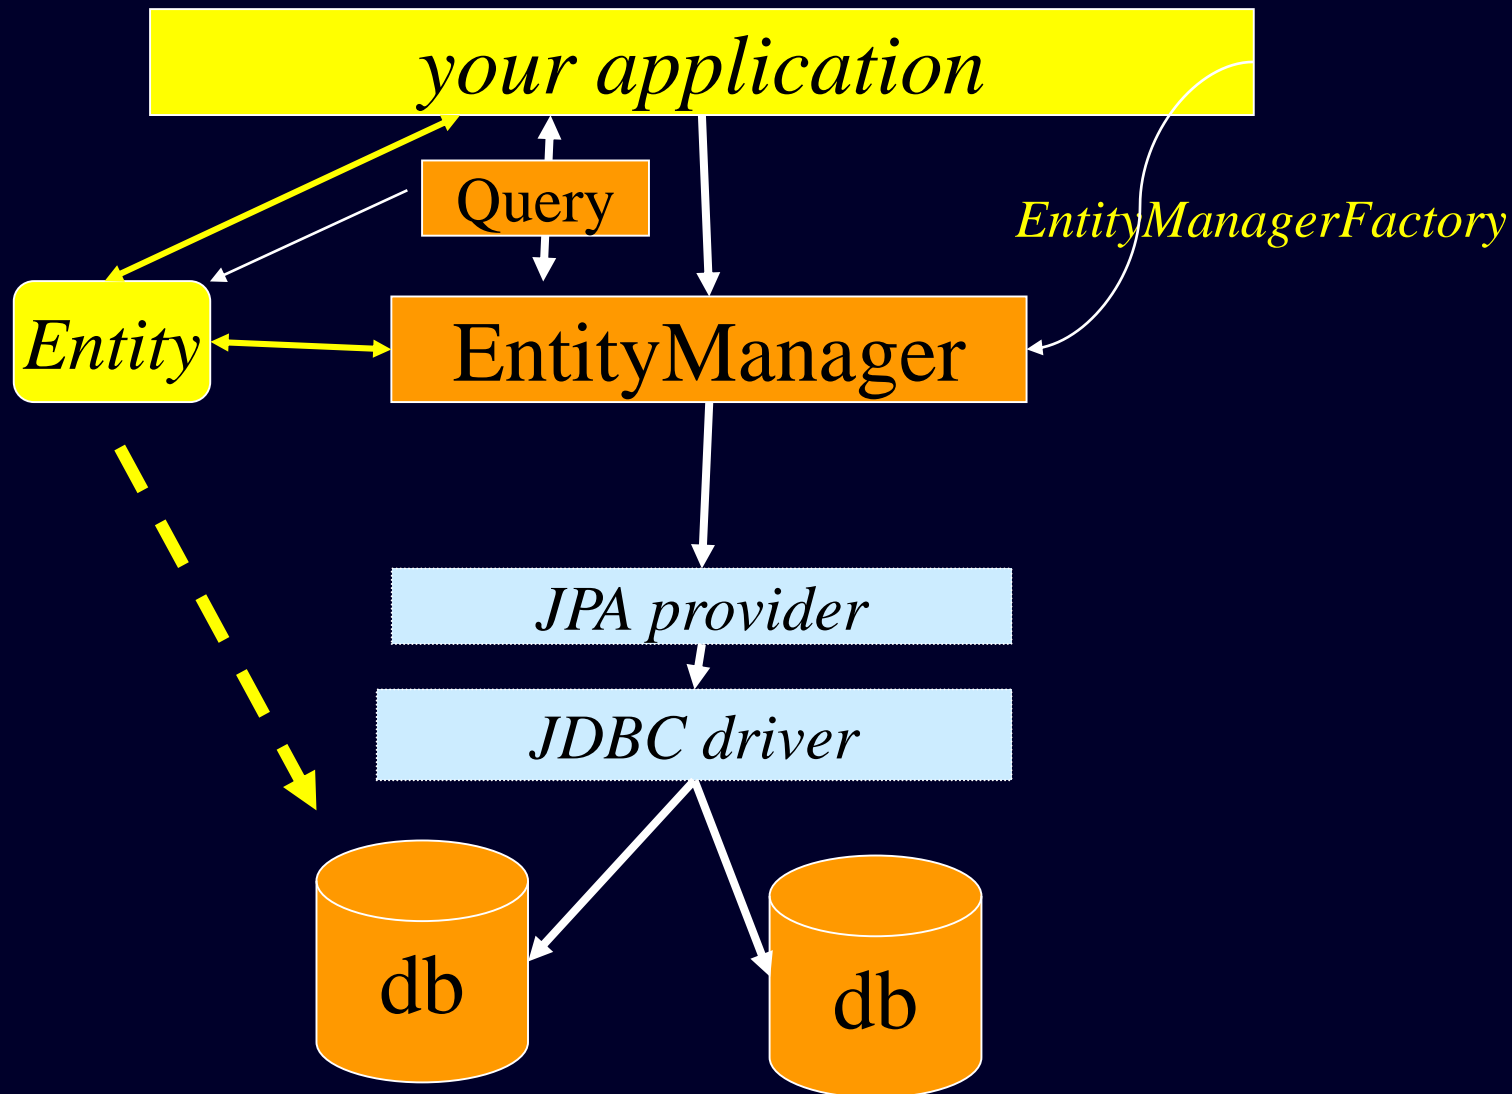
Topics

- Introduction to java persistence
 - **The Java Persistence API**
 - **Entities**
 - EntityManager & the Persistent Context
 - Persistence Units
 - Exceptions
 - JPA Query Language
-

Java Persistence API

- javax.persistence.*
 - EntityManager
 - EntityManagerFactory
 - EntityTransaction
 - Query
 - "*Entity*" – we use Plain Old Java Objects (POJO) instead.
-

JPA classes



Entities

- An **entity** is a plain old java object (POJO)
 - The **Class** represents a **table** in a relational database.
 - **Instances** correspond to **rows**
 - Requirements:
 - annotated with the **`javax.persistence.Entity`** annotation
 - **`public`** or **`protected`**, no-argument constructor
 - the class must not be declared **`final`**
 - no methods or persistent instance variables must be declared **`final`**
-

Requirements for Entities (cont.)

- May be `serializable`, but not required
 - Only needed if passed by value (in a remote call)
- Entities may extend both entity and non-entity classes
- Non-entity classes may extend entity classes
- Persistent instance variables must be declared private, protected, or package-private
- No required business/callback interfaces

- Example:

```
@Entity
class Person{
    . . .
}
```

Persistent Fields and Properties

- The persistent state of an entity can be accessed:
 - through the entity's **instance variables**
 - through **JavaBeans-style properties** (getters/setters)
 - Supported types:
 - primitive types, String, other serializable types, enumerated types
 - other entities and/or collections of entities
 - embeddable classes
 - All fields not annotated with **@Transient** or not marked as Java transient will be persisted to the data store!
-

Primary Keys in Entities

- Each entity must have a unique object identifier (persistent identifier)

@Entity

```
public class Employee {  
    @Id private int id;  
    private String name;  
    private Date age;  
  
    public int getId() { return id; }  
    public void setId(int id) { this.id = id;  
}  
    . . .  
}
```

Primary key

Persistent Identity

- Identifier (id) in entity = primary key in database
 - Uniquely identifies entity in memory and in DB
 - Persistent identity types:
 - Simple id – single field/property
`@Id int id;`
 - Compound id – multiple fields/properties
`@Id int id;`
`@Id String name;`
 - Embedded id – single field of PK class type
`@EmbeddedId EmployeePK id;`
-

Identifier Generation

- Identifiers can be generated in the database by specifying **@GeneratedValue** on the identifier
- Four pre-defined generation strategies:
 - AUTO, IDENTITY, SEQUENCE, TABLE
- Generators may pre-exist or be generated
- Specifying strategy of AUTO indicates that the provider will choose a strategy

```
@Id @GeneratedValue(strategy=AUTO)  
private int id;
```

Customizing the Entity Object

- In most of the cases, the defaults are sufficient
- By default the table name corresponds to the unqualified name of the class
- Customization:

```
@Entity
@Table(name = "FULLTIME_EMPLOYEE")
public class Employee{ ..... }
```

- The defaults of columns can be customized using the @Column annotation

```
@Id @Column(name = "EMPLOYEE_ID", nullable = false)
private String id;

@Column(name = "FULL_NAME" nullable = true, length = 100)
private String name;
```

Entity Relationships

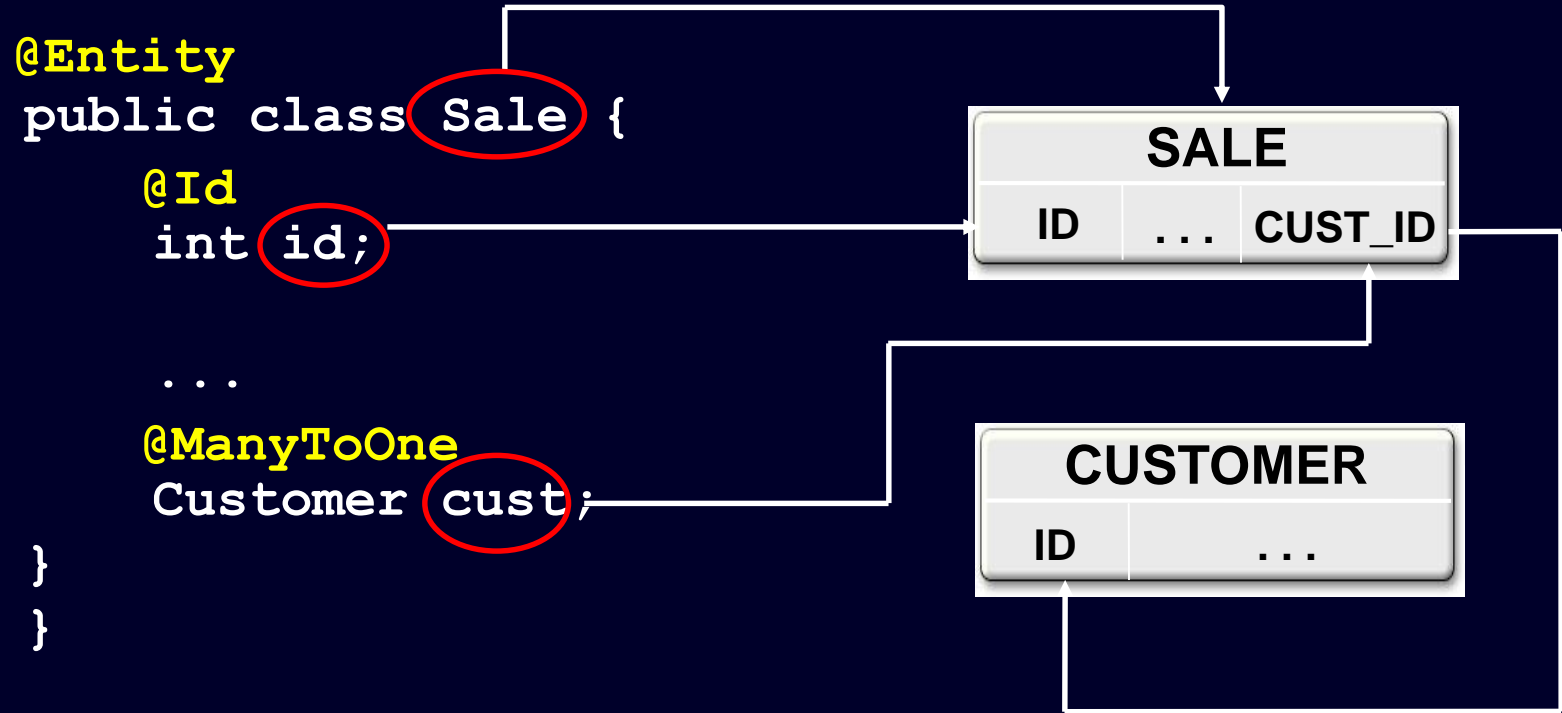
- There are four types of relationship multiplicities:
 - @OneToOne
 - @OneToMany
 - @ManyToOne
 - @ManyToMany
 - The direction of a relationship can be:
 - **bidirectional** – owning side and inverse side
 - **unidirectional** – owning side only
 - Owning side specifies the physical mapping
-

Entity Relation Attributes

- JPA supports cascading updates/deletes
 - CascadeType
 - ALL, PERSIST, MERGE, REMOVE, REFRESH
- You can declare performance strategy to use with fetching related rows
 - FetchType
 - LAZY, EAGER
 - (Lazy means don't load row until the property is retrieved)

```
@ManyToMany(  
    cascade = {CascadeType.PERSIST, CascadeType.MERGE},  
    fetch = FetchType.EAGER)
```

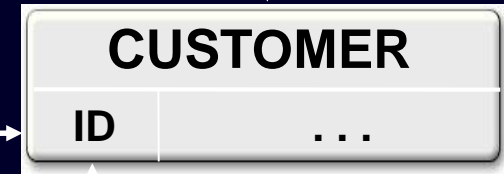
ManyToOne Mapping



OneToMany Mapping

```
@Entity
public class Customer {
    @Id
    int id;
    ...
    @OneToMany(mappedBy="cust")
    Set<Sale> sales;
}

@Entity
public class Sale {
    @Id
    int id;
    ...
    @ManyToOne
    Customer cust;
}
```



ManyToOne Mapping

```
@Entity
public class Customer {
    ...
    @JoinTable(
        name="CUSTOMER_SALE",
        joinColumns=@JoinColumn(
name="CUSTOMER_ID",referencedColumnName="customer_id"),
        inverseJoinColumns=@JoinColumn(
name="SALE_ID", referencesColumnName="sale_id")
        Collection<Sale> sales;
    }
```

```
@Entity
public class Sale {
    ...
    @ManyToOne(mappedBy="sales")
    Collection<Customer> customers;
}
```

Entity Inheritance

- An important capability of the JPA is its support for inheritance and polymorphism
 - Entities can inherit from other entities and from non-entities
 - The **@Inheritance** annotation identifies a mapping strategy:
 - SINGLE_TABLE
 - JOINED
 - TABLE_PER_CLASS
-

Inheritance Example

```
@Entity
@Inheritance(strategy=InheritanceType.SINGLE_TABLE)
@DiscriminatorColumn(name="DISC",
                    discriminatorType=DiscriminatorType.STRING)
@DiscriminatorValue(name="CUSTOMER")
public class Customer { . . . }

@Entity
@DiscriminatorValue(name="VCUSTOMER")
public class ValuedCustomer extends Customer { . . . }
```

- SINGLE_TABLE strategy - all classes in the hierarchy are mapped to a single table in the database
 - Discriminator column - contains a value that identifies the subclass
 - Discriminator type - {STRING, CHAR, INTEGER}
 - Discriminator value - value entered into the discriminator column for each entity in a class hierarchy
-

Topics

- Introduction to java persistence
 - The Java Persistence API
 - Entities
 - **EntityManager & the Persistent Context**
 - Persistence Units
 - Exceptions
 - JPA Query Language
-

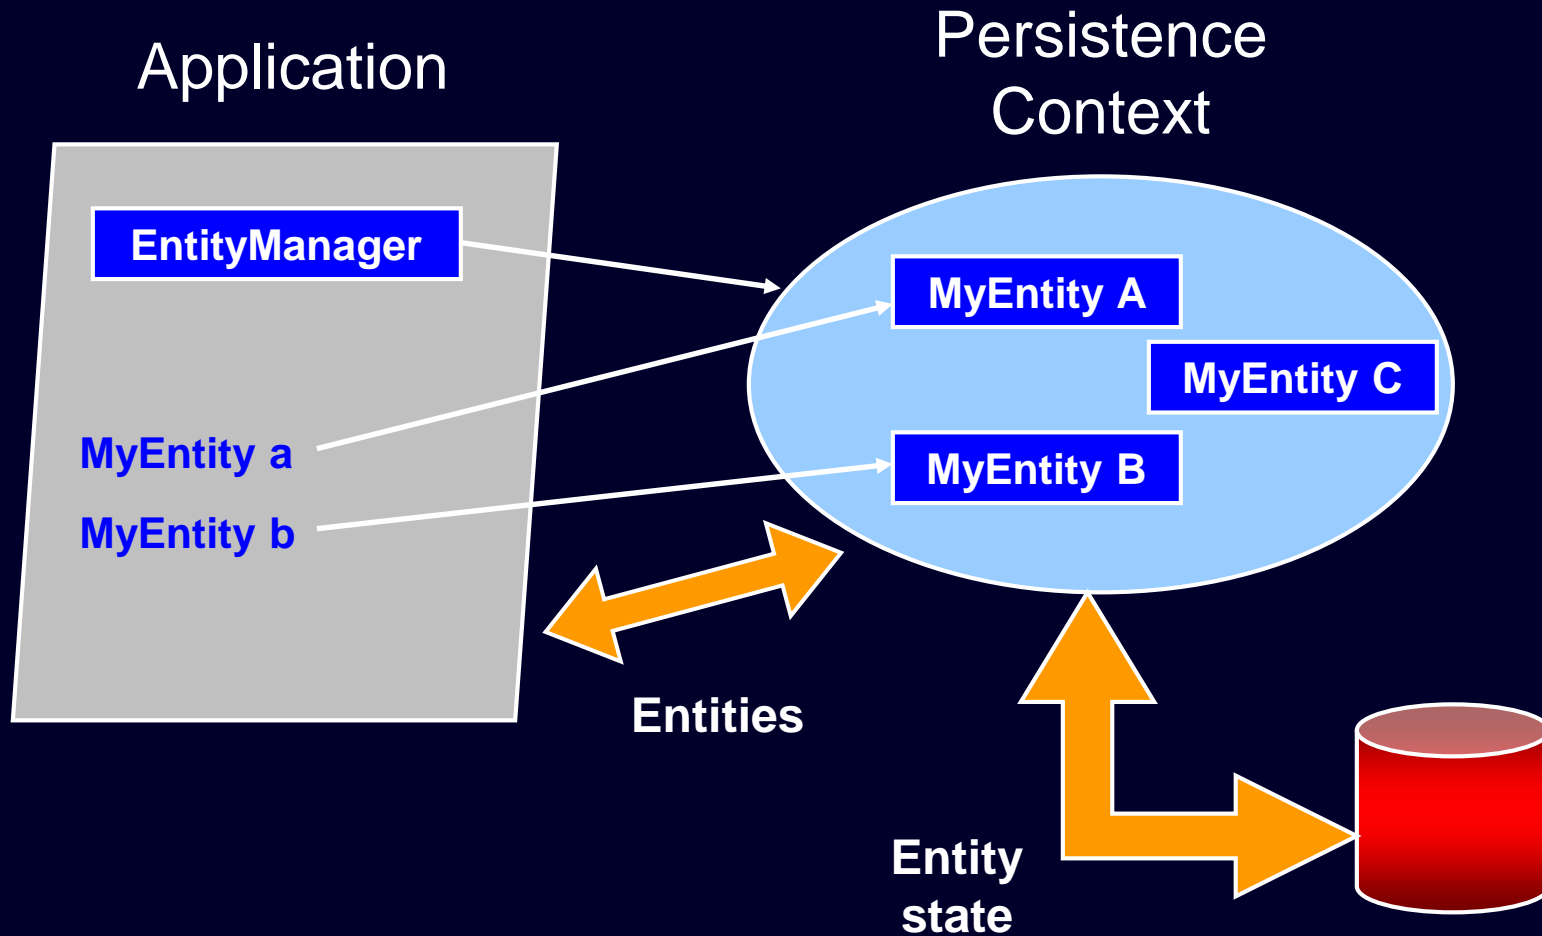
Managing Entities

- Entities are managed by the **entity manager**
 - The entity manager is represented by `javax.persistence.EntityManager` instances
 - Each EntityManager instance is associated with a **persistence context**
 - A persistence context defines the scope under which particular entity instances are created, persisted, and removed
-

Persistence Context

- A **persistence context** is a set of managed entity instances that exist in a particular data store
 - Entities **keyed** by their persistent **identity**
 - Only **one** entity with a given persistent identity may exist in the persistence context
 - Entities are added to the persistence context, but are not individually removable (“detached”)
 - Controlled and managed by **EntityManager**
 - Contents of persistence context change as a result of operations on EntityManager API
-

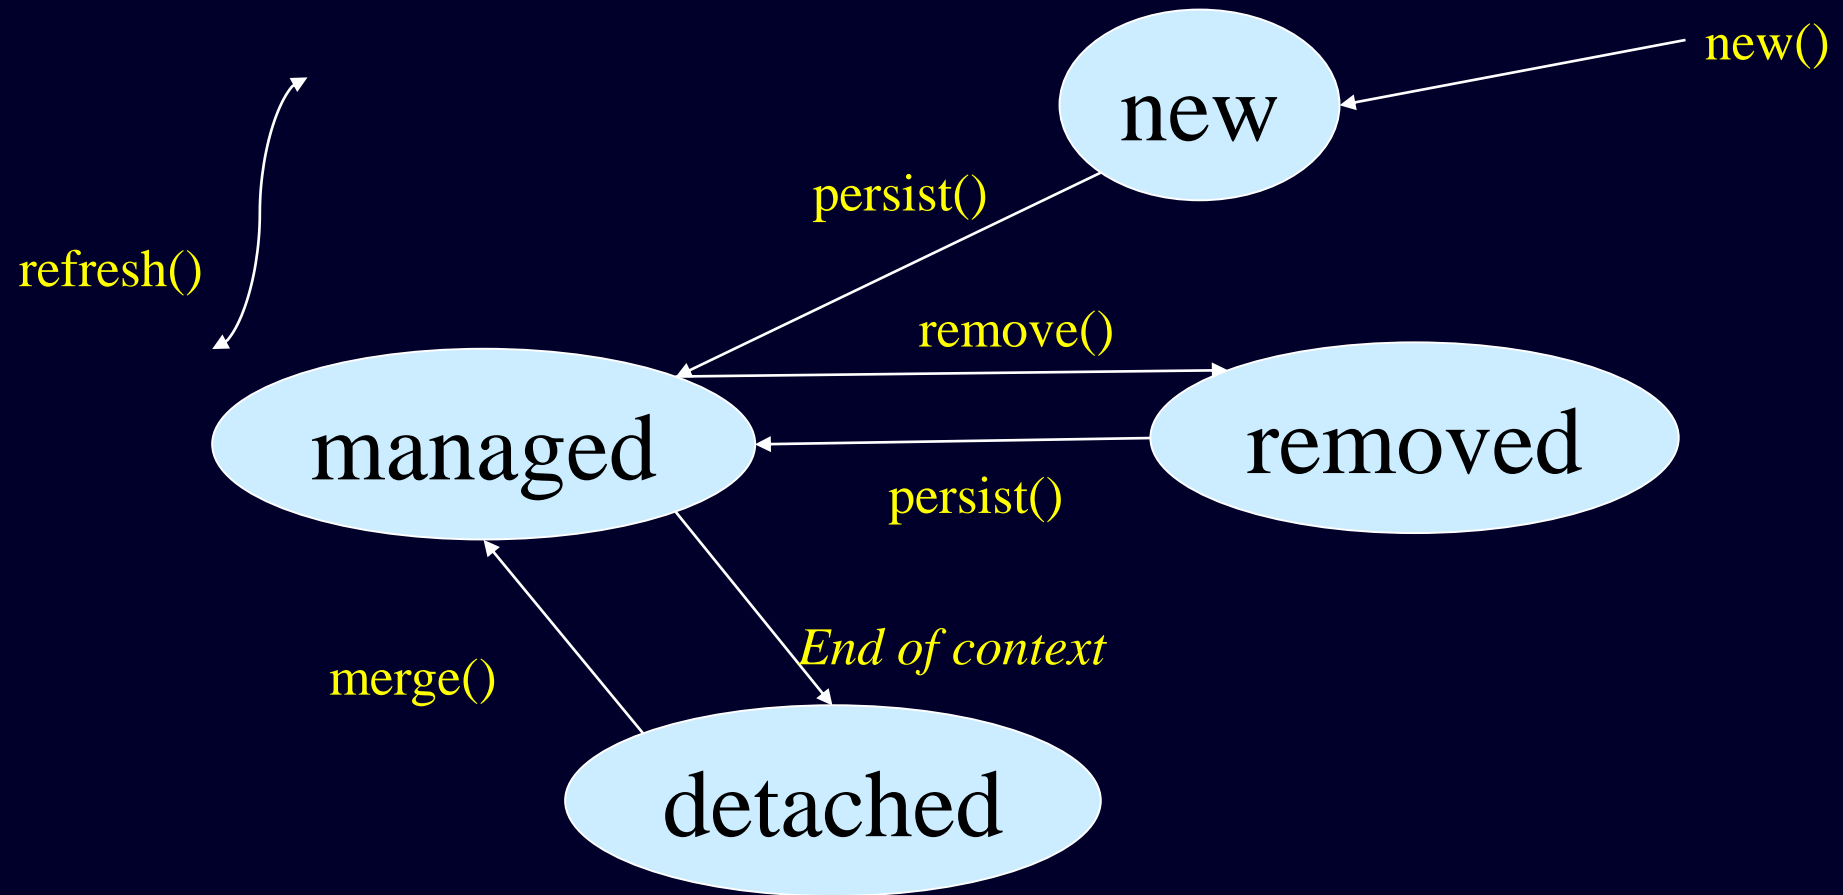
Persistence Context



Entity Manager

- An **EntityManager** instance is used to manage the state and life cycle of entities within a persistence context
 - Entities can be in one of the following states:
 1. New
 2. Managed
 3. Detached
 4. Removed
-

Entity Lifecycle



Entity Lifecycle

- **New** – entity is instantiated but not associated with persistence context. Not linked to database.
 - **Managed** – associated with persistence context. Changes get synchronised with database
 - **Detached** – has an id, but not connected to database
 - **Removed** – associated with persistence context, but underlying row will be deleted.
 - The **state** of persistent entities is synchronized to the database when the transaction **commits**
-

Entity Manager

- The EntityManager API:
 - **creates** and **removes** persistent entity instances
 - **finds** entities by the entity's primary key
 - allows **queries** to be run on entities
 - There are two types of EntityManagers:
 - **Application-Managed** EntityManagers
 - ie: run via Java SE
 - **Container-Managed** EntityManagers
 - ie: run via Java EE Container eg: Tomcat
-

Java SE applications create EntityManager instances by using **directly** Persistence and EntityManagerFactory:

- **javax.persistence.Persistence**
 - Root class for obtaining an EntityManager
 - Locates provider service for a named persistence unit
 - Invokes on the provider to obtain an EntityManagerFactory
 - **javax.persistence.EntityManagerFactory**
 - Creates EntityManagers for a named persistence unit or configuration
-

- Applications must manage own transactions too..

```
public class PersistenceProgram {  
    public static void main(String[] args)  
    {  
        EntityManagerFactory emf =  
Persistence.createEntityManagerFactory("SomePUnit");  
        EntityManager em = emf.createEntityManager();  
        em.getTransaction().begin();  
        // Perform finds, execute queries,  
...  
        // update entities, etc.  
        em.getTransaction().commit();  
        em.close();  
        emf.close();  
    }  
}
```

- Containers provide naming and transaction services for JPA
 - (eg: Web Container like Tomcat, EJB Container like WebLogic)
 - JPA relies on the container to insert the actual reference to the EntityManager for the current context via ***dependency injection***
 - Use the Java5 annotations to do this
-

Container-Managed EntityManager

- An EntityManager with a transactional persistence context can be **injected** by using the **@PersistenceContext** annotation

```
@PersistenceContext (unitName="SomePUnit")
    EntityManager em;
// Perform finds, execute queries,
...
// update entities, etc.
em.close();
```

container inserts
reference to the
container's
EntityManager

- You could also use the **@Resource(name="jndi:name")** annotation to insert a named entity manager.

Transactions

- JPA transactions can be managed by:
 - the users application
 - a framework (such as Spring)
 - a Java EE container
 - Transactions can be controller in two ways:
 - **Java Transaction API (JTA)**
 - container-managed entity manager
 - **EntityManager API** (`tx.begin()`, `tx.commit()`, etc)
 - application-managed entity manager
-

Operations on Entity Objects

- EntityManager API operations:
 - **persist()** - Save the entity into the db
 - **remove()** - Delete the entity from the db
 - **refresh()** - Reload the entity state from the db
 - **merge()** - Synchronize a detached entity with the p/c
 - **find()** - Find by primary key
 - **createQuery()** - Create query using dynamic JP QL
 - **createNamedQuery()** - Create a predefined query
 - **createNativeQuery()** - Create a native "pure" SQL query.
Can also call stored procedures.
 - **contains()** - Is entity is managed by p/c
 - **flush()** - Force synchronization of p/c to database

Note: p/c == the current persistence context

Topics

- Introduction to java persistence
 - The Java Persistence API
 - Entities
 - EntityManager & the Persistent Context
 - **Persistence Units**
 - Exceptions
 - JPA Query Language
-

Persistence Units

- A **persistence unit** defines a set of all entity classes that are managed by `EntityManager` instances in an application
 - Each persistence unit can have different providers and database drivers
 - Persistence units are defined by the **persistence.xml** configuration file
-

persistence.xml

A **persistence.xml** file defines one or more persistence units

```
<persistence>
  <persistence-unit name="SomePUnit">
    <provider>org.hibernate.ejb.HibernatePersistence</provider>
    <class>myapp.MyEntity</class>
    <properties>
      <property name="hibernate.connection.url"
        value="jdbc:oracle:thin:@smaug.it.uts.edu.au:1522:ell"/>
      <property name="hibernate.connection.driver_class"
        value="oracle.jdbc.driver.OracleDriver"/>
      <property name="hibernate.connection.username"
        value="user"/>
      <property name="hibernate.connection.password"
        value="password"/>
    </properties>
  </persistence-unit>
</persistence>
```

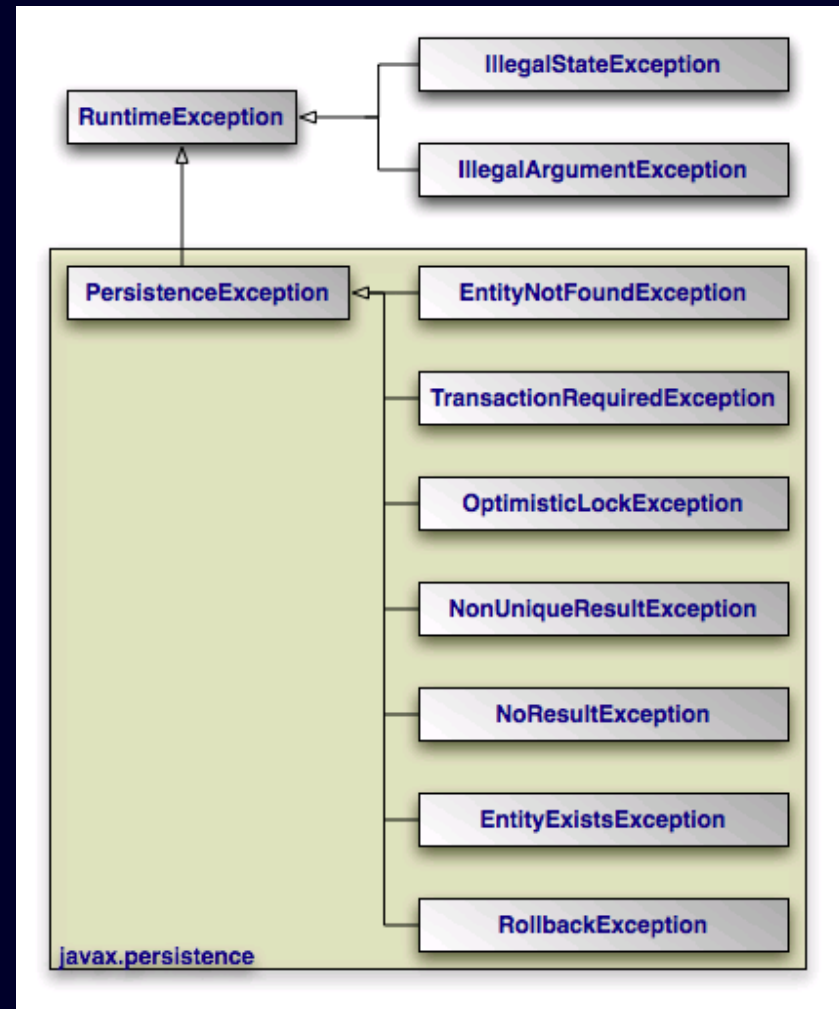
persistence.xml

You can also use JNDI + datasource in `persistence.xml` instead of hard coding driver details. Requires container to manage this.

```
<persistence>
  <persistence-unit name=" SomePUnit">
    <provider>org.hibernate.ejb.HibernatePersistence</provider>
    <jta-data-source>jdbc/thinOracleDataSource</jta-data-source>
    <class>myapp.MyEntity</class>
  </persistence-unit>
</persistence>
```

JPA exceptions

- All exceptions are unchecked
- Exceptions in `javax.persistence` package are self-explanatory



Topics

- Introduction to java persistence
 - The Java Persistence API
 - Entities
 - EntityManager & the Persistent Context
 - Persistence Units
 - Exceptions
 - **JPA Query Language**
-

JPQL Introduction

- JPA has a query language based on SQL
 - JPQL is an extension of EJB QL
 - More robust flexible and object-oriented than SQL
 - The persistence engine parses the query string, transform the JPQL to the native SQL before executing it
-

Creating Queries

- Query instances are obtained using:
 - EntityManager.**createNamedQuery** (static query)
 - EntityManager.**createQuery** (dynamic query)
 - EntityManager.**createNativeQuery** (native query)
 - Query API:
 - **getResultList()** – execute query returning multiple results
 - **getSingleResult()** – execute query returning **single** result
 - **executeUpdate()** – execute bulk update or delete
 - **setFirstResult()** – set the first result to retrieve
 - **setMaxResults()** – set the maximum number of results to retrieve
 - **setParameter()** – bind a value to a named or positional parameter
 - **setHint()** – apply a vendor-specific hint to the query
 - **setFlushMode()** – apply a flush mode to the query when it gets run
-

Static (Named) Queries

- Defined statically with the help of **@NamedQuery** annotation together with the entity class
- @NamedQuery elements:
 - **name** - the name of the query that will be used with the createNamedQuery method
 - **query** – query string

```
@NamedQuery(name="findAllCustomers",  
            query="SELECT c FROM Customer")
```

```
Query findAllQuery =  
entityManager.createNamedQuery("findAllCustomers");  
List customers = findAllQuery.getResultList();
```

Multiple Named Queries

Multiple named queries can be logically defined with the help of **@NamedQueries** annotation

```
@NamedQueries( {  
    @NamedQuery(name = "Mobile.selectAllQuery"  
                query = "SELECT M FROM MOBILEENTITY"),  
    @NamedQuery(name = "Mobile.deleteAllQuery"  
                query = "DELETE M FROM MOBILEENTITY")  
} )
```

Dynamic Queries

- Dynamic queries are queries that are defined directly within an application's business logic
- ! Not efficient & slower. Persistence engine has to parse, validate & map the JPQL to SQL at run-time

```
public List findAll(String entityName) {  
    return entityManager.createQuery(  
        "select e from " + entityName + " e")  
        .getResultList();  
}
```

Named Parameters

- Named parameters are parameters in a query that are prefixed with a **colon** (:)
- To bound parameter to an argument use method:
 - **Query.setParameter**(String name, Object value)

```
public List findWithName(String name) {  
    return em.createQuery(  
        "SELECT c FROM Customer c WHERE c.name LIKE :custName")  
        .setParameter("custName", name)  
        .getResultList();  
}
```

Positional Parameters

- Positional parameters are prefixed with a **question mark (?)** & number of the parameter in the query
- To set parameter values use method:
 - `Query.setParameter(integer position, Object value)`

```
public List findWithName(String name) {  
    return em.createQuery(  
        "SELECT c FROM Customer c WHERE c.name LIKE ?1")  
        .setParameter(1, name)  
        .getResultList();  
}
```

Native Queries

- Queries may be expressed in native SQL
- Use when you need to use native SQL of the target database
- Can call stored procedures using "call procname" syntax

```
Query q = em.createNativeQuery(  
    "SELECT o.id, o.quantity, o.item " +  
    "FROM Order o, Item i " +  
    "WHERE (o.item = i.id) AND (i.name = 'widget')",  
    com.acme.Order.class);
```

- Use **@SqlResultSetMapping** annotation for more advanced cases
-

Query Operations – Multiple Results

- `Query.getResultList()` will execute a query and may return a List object containing multiple entity instances

```
Query query = entityManager.createQuery("SELECT C FROM CUSTOMER");  
List<MobileEntity> mobiles = (List<MobileEntity>) query.getResultList();
```

- Will return a non-parameterized List object
 - Can only execute on select statements as opposed to UPDATE or DELETE statements
 - For a statement other than SELECT run-time `IllegalStateException` will be thrown
-

Query Operations – Single Result

- A query that returns a single entity object

```
Query singleSelectQuery = entityManager.createQuery(  
    "SELECT C FROM CUSTOMER WHERE C.ID = 'ABC-123'");  
Customer custObj = singleSelectQuery.getSingleResult();
```

- If the match wasn't successful, then **EntityNotFoundException** is returned
 - If more than one matches occur during query execution a run-time exception **NonUniqueResultException** will be thrown
-

Paging Query Results

```
int maxRecords = 10; int startPosition = 0;
String queryString = "SELECT M FROM MOBILEENTITY";
while(true){
    Query selectQuery = entityManager.createQuery(queryString);
    selectQuery.setMaxResults(maxRecords);
    selectQuery.setFirstResult(startPosition);
    List<MobileEntity> mobiles =

        entityManager.getResultList(queryString);
    if (mobiles.isEmpty()){ break; }
    process(mobiles);           // process the mobile entities
    entityManager.clear();      // detach the mobile objects
    startPosition = startPosition + mobiles.size();
}
```

Flushing Query Objects

- Two modes of flushing query objects
 - **AUTO** (default) and **COMMIT**
 - **AUTO** - any changes made to entity objects will be reflected the very next time when a SELECT query is made
 - **COMMIT** - the persistence engine may only update all the state of the entities during the database COMMIT
 - set via `Query.setFlushMode()`
-

JPQL Statement Language

- JPQL statement types:
 - SELECT, UPDATE, DELETE
 - Supported clauses:
 - FROM
 - WHERE
 - GROUP_BY
 - HAVING
 - ORDER BY
 - ...
 - Conditional expressions, aggregate functions,...
-

JPQL Enhancements over EJBQL 2.x

- Simplified query syntax
 - JOIN operations
 - Group By and Having Clause
 - Subqueries
 - Dynamic queries
 - Named parameters
 - Bulk update and delete
-

OO-style vs. SQL-style queries

- The main difference:
*** query the application model, i.e. the entities,
rather than any database tables
- Better productivity by using OO-style queries, e.g.

`employee.getManager().getAddress()`

which becomes:

```
SELECT t3.* FROM EMP t1, EMP t2, ADDR t3
WHERE t1.EMP_ID = "XYZ" AND t1.MGR_ID = t2.EMP_ID
AND t2.ADDR_ID = t3.ADDR_ID
```

- Notice that the two-step object traversal was packed
into a single DB query
-

Questions?



- The Java Persistence API - A Simpler Programming Model for Entity Persistence
<http://java.sun.com/developer/technicalArticles/J2EE/jpa/>
 - Article "Introduction to Java Persistence API"
http://www.javabeat.net/javabeat/ejb3/articles/2007/04/introduction_to_java_persistence_api_jpa_ejb_3_0_1.php
 - TopLink Essentials (reference implementation)
<https://glassfish.dev.java.net/javaee5/persistence/>
 - JPA Annotation Reference
<http://www.oracle.com/technology/products/ias/toplink/jpa/resources/toplink-jpa-annotations.html>
-

- JPQL Language Reference

http://openjpa.apache.org/builds/1.0.2/apache-openjpa-1.0.2/docs/manual/jpa_langref.html

- JPA Query API

http://www.javabeat.net/javabeat/ejb3/articles/2007/04/introduction_to_java_persistence_api_jpa_ejb_3_0_6

- Standardizing Java Persistence with the EJB3 Java Persistence API – Query API

<http://www.onjava.com/pub/a/onjava/2006/05/17/standardizing-with-ejb3-java-persistence-api.html?page=last&x-showcontent=text>